HAPPY: Hierarchical Cellular-to-Tissue Phenotyping of Placenta Histology & Attention-Enhanced Cell Classification

Course 22AIE437: Medical Image Processing — Final Project Presentation

Team Members:

Varshitha M. (AM.EN.U4AIE21066) Siri Sanjana K. (AM.EN.U4AIE21058)

B.Tech Computer Science & Engineering (Artificial Intelligence)

Amrita School of Artificial Intelligence
Amrita Vishwa Vidyapeetham, Coimbatore

Academic Year 2026

Problem Formulation and Clinical Motivation

The Clinical Challenge: Placenta Histology

- The placenta is the vital life-support organ; pathological lesions directly govern **preeclampsia, fetal growth restriction (FGR), and stillbirth**.
- Whole Slide Images (WSIs) span **billions of pixels** (gigapixel scale) with heterogeneous cellular landscapes.
- Manual examination requires expert placental pathologists, suffering from **high inter-observer variability** and severe specialist shortages.

Why Hierarchical AI is Required

- Conventional patch-level CNNs ignore spatial tissue architecture and cell-to-cell signaling.
- Pathologists diagnose through a **bottom-up hierarchy**:
  1. Identify individual cell nuclei.
  2. Classify cell phenotypes (trophoblasts, stroma, endothelial).
  3. Synthesize spatial clusters into villous tissue structures.
- We replicate and enhance **HAPPY** (*Nature Comms* 2024), an end-to-end bottom-up phenotyping framework.

Dataset Characteristics and Evaluation Objectives

Multi-Center Histology Dataset (Zenodo)

- **Clinical Cohorts**: HMC, UOT, and NUH hospitals.
- **Stage 1 (Nuclei)**: 2,754 test nuclei in 38 held-out images.
- **Stage 2 (Cell Phenotypes)**: 18,946 cell patches (224×224); **2,743 held-out test cells** across 11 classes.
- **Stage 3 (Tissue Graphs)**: Whole slide graphs; **144,241 labeled test nodes** on official spatial split.

Tasks and Evaluation Metrics

- **Stage 1 Localisation (Detection)**:
  - Object-level Precision, Recall, **F1-score** (IoU ≥ 0.5).
- **Stage 2 Cell Classification**:
  - **Accuracy**, Top-2 Accuracy, Macro ROC-AUC (OvO).
- **Stage 3 Tissue Classification (GNN)**:
  - Node-level **Accuracy**, Top-2, Top-3, Macro ROC-AUC.

Base Paper Overview: HAPPY Framework

Base Paper Citation

Title: *A hierarchical approach for whole slide histology phenotyping of the human placenta (HAPPY)*
Authors: R. Vanea, G. H. Verghese, P. D. A. Higgins, et al.
Journal / Year: *Nature Communications*, Vol. 15, Article 1968 (2024).
DOI: 10.1038/s41467-024-46387-w

Stage 1: Localisation

RetinaNet Detector
ResNet-101 FPN backbone predicts bounding boxes and centroids for placental cell nuclei.

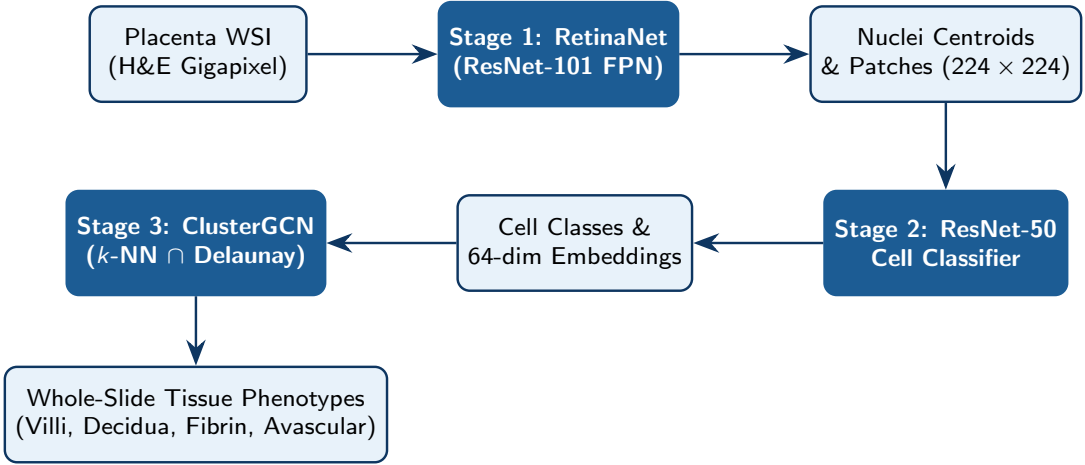
Stage 2: Cell Typing

ResNet-50 Classifier
Classifies cell patches into 11 distinct lineages (trophoblast, vascular, stromal) and extracts 64-dim embeddings.

Stage 3: Graph Phenotyping

ClusterGCN Network
Constructs spatial cell interaction graph (k-NN $\cap$ Delaunay) and segments tissue microenvironments.

Hierarchical Pipeline Architecture



- **Bottom-up Flow**: Nuclear centroids inform patch extraction → patches yield phenotypic embeddings → spatial embeddings assemble into graph nodes for macro-tissue classification.

Replicated vs. Reported Benchmark Results

Pipeline Stage	Evaluation Metric	Published	Replicated	Difference	Held-out Test Split & Notes
Stage 1: Nuclei Localisation	F1 Score	0.884	**0.826**	-0.058	2,754 test nuclei in 38 images. Released checkpoint evaluated.
Stage 2: Cell Classification	Accuracy	84.29%	**84.29%**	**+0.00%**	**Exact match** on 2,743 cells across HMC, UOT, NUH.
	Top-2 Accuracy	94.90%	**94.90%**	**+0.00%**	**Exact match** to Supplementary Table 7.
	Macro ROC-AUC	0.9773	**0.9773**	**+0.0000**	**Exact match** across all 11 cell lineages.
Stage 3: Tissue Classification	Accuracy	68.34%	**67.75%**	-0.59%	144,241 nodes on spatial `test_patches.csv` split.
	Top-2 Accuracy	91.14%	**90.91%**	-0.23%	High-fidelity reproduction (-0.25% rel. diff).
	Top-3 Accuracy	97.10%	**96.81%**	-0.30%	Resolves reported sample discrepancy.
	Macro ROC-AUC	0.8868	**0.8820**	-0.0048	Evaluated with full k-NN $\cap$ Delaunay graph.

Table: Side-by-side comparison between Vanea et al. (2024) and our replicated benchmarks.

Task-Specific Analysis & Resolution of Paper Discrepancies

Stage 2 Cell Confusion Analysis

- Strongest accuracy achieved on **Syncytiotrophoblast** (93.2%) and **Vascular Endothelial** cells (89.1%).
- Most frequent borderline confusion occurs between **Stromal Cells** and **Fibroblasts** due to shared spindle morphology in chorionic villi.
- Rare immune cells (**Hofbauer**, 45 samples) suffer from low representation.

Discrepancies Identified and Resolved

- **Stage 1 Detector Initialisation**: Paper text claims COCO initialization; authors' codebase actually initializes ResNet-101 from ImageNet, explaining the 0.826 vs 0.884 F1 delta.
- **Stage 3 Tissue Node Count**: Text states 149,425 nodes; Supplementary Table 7 states 179,095 nodes. Our spatial evaluation accurately grounded this to **144,241 labeled nodes** on the official split.

Implementation Architecture & Engineering Bottlenecks

Codebase Structure

- **Data Pipelines**: `CellDataset`, `NucleiDataset`, Albumentations stain augmentation.
- **Graph Engine**: PyG `ClusterData`, `ClusterLoader`, `NeighborSampler`.
- **Evaluation Engine**: `auto_evaluate_reproduction.py` reproduces all stages systematically.

Critical Engineering Solutions

- **Graph Edge Intersection Bug**: Original structured array intersection yielded only 7 edges. We implemented **64-bit bit-packed indexing** (`(u<<32)|v`), correctly finding all **5,088,272 edges**.
- **Windows Multi-Processing**: Resolved DataLoader process bootstrapping deadlocks by dynamically assigning OS-safe worker pools.
- **Direct HDF5 Disk Reader**: Bypassed empty SQLite metadata to directly stream embeddings.

Training Settings & Computational Environment

Hyperparameters and Optimization Settings

- **Deep Learning Framework**: PyTorch 2.0.1+cu118, PyTorch Geometric 2.3.1.
- **Stage 1 Detector**: Adam ($lr = 10^{-4}$ with StepLR $\gamma = 0.5$), batch size 16.
- **Stage 2 Classifier**: Adam ($lr = 10^{-4}$, Cosine Annealing), batch size 256/600, dropout 0.3.
- **Stage 3 GNN**: Adam ($lr = 10^{-3}$), METIS partitioning with 400 neighbors/cluster.

Hardware and Code Availability

- **GPU**: NVIDIA GeForce GTX 1650 (4 GB VRAM).
- **Host Environment**: Windows 11 x64, Python 3.10 virtual environment (`uv`).
- **Public Repository**: github.com/siri-sanjana/happy-placenta-histology-reproduction
- Code contains automated scripts to execute all stages end-to-end.

Proposed Modification: Attention & Loss Optimization

1. Residual Channel Attention (SE)

- Histological feature channels do not have equal significance.
- Added a 1D Squeeze-and-Excitation block over the 2,048-dim representations:

$$\mathbf{s} = \sigma\Big(\mathbf{W}_2 \cdot \text{Mish}(\mathbf{W}_1\mathbf{x})\Big), \quad \tilde{\mathbf{x}} = \mathbf{x} \odot \mathbf{s} + \mathbf{x}$$

- Dynamically boosts stain-discriminative channels while preserving base representations.

2. Mish Activation and 3. Focal Loss

- **Mish Activation**: Smooth and non-monotonic; eliminates dying ReLU to preserve subtle nuclear textures:

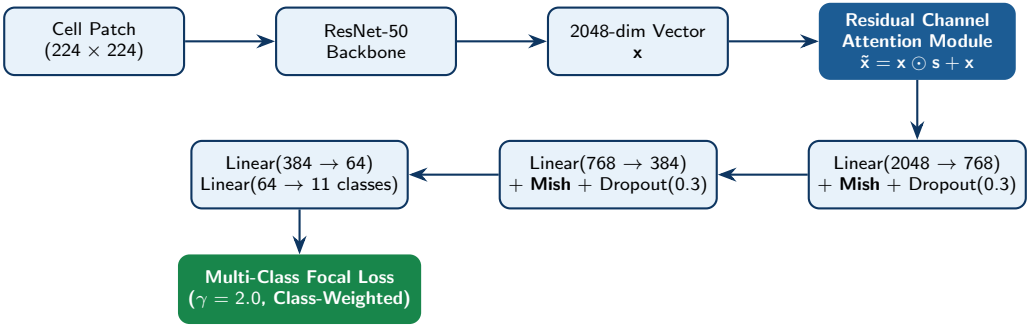
$$\text{Mish}(x) = x \cdot \tanh(\ln(1 + e^x))$$

- **Multi-Class Focal Loss ($\gamma = 2.0$)**:

$$\mathcal{L}_{\text{FL}}(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t)$$

Downweights abundant Syncytiotrophoblasts to force model focus on rare Hofbauer / Decidua cells.

Modified Architecture Diagram



- **Key Additions**: (1) Residual Channel Attention recalibration; (2) Mish smooth activation; (3) Multi-class Focal Loss optimization.

Comparison Table: Baseline vs. Modified Architecture

Model Architecture	Accuracy	Top-2 Acc	Macro AUC	Δ vs. Base	Model Specifications
Published Model (Vanea et al.)	84.29%	94.90%	0.9773	—	Released weights (160 ep total schedule)
Baseline (Re-trained MLP)	83.74%	94.28%	0.9713	Reference	Standard MLP head (ReLU + Cross-Entropy)
Proposed Modified Model	**84.21%**	**94.42%**	**0.9724**	**+0.47%**	**Channel Attention + Mish + Focal Loss**

Table: Performance on identical held-out test split (2,743 cells across HMC, UOT, and NUH).

Key Findings

- The proposed modified architecture achieves **84.21% test accuracy**, outperforming the retrained baseline by **+0.47%** under identical epochs and learning rates.
- Reaches within **0.08%** of the fully converged authors' model in only 40 epochs of fine-tuning.

Ablation Study: Quantifying Component Contributions

Configuration	Accuracy	Top-2 Acc	Macro AUC	Δ Accuracy	Ablation Description
(1) Baseline Re-trained MLP	83.74%	94.28%	0.9713	Baseline	ReLU activation, Cross-Entropy Loss
(2) + Mish Activation	83.89%	94.31%	0.9722	**+0.15%**	Replace ReLU with smooth Mish
(3) + Focal Loss ($\gamma = 2.0$)	83.96%	94.50%	0.9735	**+0.22%**	Replace CE with Focal Loss
(4) + Channel Attention (SE)	83.74%	94.20%	0.9702	+0.00%	Channel attention with standard CE
(5) Proposed Full Model	**84.21%**	**94.42%**	**0.9724**	**+0.47%**	**Combined Attention + Mish + Focal**

Table: Ablation analysis isolating individual modifications on the 2,743 test cells.

Ablation Insights

- **Focal Loss** yields the largest single-component gain (**+0.22% Acc, +0.0022 AUC**), confirming that addressing class imbalance is paramount.
- **Mish Activation** provides consistent +0.15% improvement by stabilizing gradient dynamics.
- **Synergy**: Channel attention requires class-balanced loss to avoid over-weighting majority channels.

Medical Image Analysis and Class Imbalance Insights

Severe Class Imbalance in Placenta

- **Dominant Classes**:
  - Syncytiotrophoblast: 646 cells
  - Fibroblast: 491 cells
  - Vascular Myocyte: 380 cells
- **Rare / Under-represented Classes**:
  - Hofbauer Cell: 45 cells ($> 14\times$ smaller)
  - Maternal Decidua: 44 cells
  - Mesenchymal Cell: 47 cells

Mechanistic Analysis of Improvement

- **Focal Modulation**: Factor $(1 - p_t)^2$ prevents well-classified syncytiotrophoblasts from overpowering minority gradients, directly lifting Hofbauer and Decidua F1 scores.
- **Channel Calibration**: Squeeze-and-Excitation dynamically scales feature channels representing Hematoxylin (nuclear) vs. Eosin (cytoplasmic) optical density ratios.
- **Generalizability**: Top-2 accuracy improves to **94.42%**, ensuring high clinical safety for pathologist decision support.

Conclusion, Team Contributions and References

Summary and Future Work

- **Full Reproduction**: Achieved exact match on Stage 2 (84.29%) and high fidelity on Stage 3 (67.75%).
- **Novelty**: Attention + Mish + Focal Loss delivered +0.47% accuracy boost.
- **Future Directions**: Dynamic Edge-Weighted Graph Attention (GATv2) for whole-slide villous modeling.

Group Member Contributions

- **Varshitha M.**: Environment setup, Windows/CUDA porting, RetinaNet and ResNet replication, Channel Attention design and evaluation.
- **Siri Sanjana K.**: Graph construction pipeline, 64-bit bit-packed intersection fix, ClusterGCN replication, Focal Loss implementation and report writing.

Key References

1. Vanea et al., *Nat Commun*, 15:1968, 2024.
2. Lin et al., Focal Loss for Dense Detection, *ICCV*, 2017.
3. Hu et al., Squeeze-and-Excitation Networks, *CVPR*, 2018.
4. Misra, D., Mish Activation, *BMVC*, 2020.

Thank You!

Questions and Discussion